

NAME: HRISHIKESH SHIVPUTRA KAMBLE

PROJECT:-

SOLVING CLASSIFICATION PREDICTION FOR "SMART HOME DEVICE USAGE" DATASET USING "LOGISTIC REGRESSION", "SVC", "GAUSSIAN NB", "BERNOULLI NB", "COMPLEMENT NB", "MULTINOMIAL NB", "K NEAREST NEIGHBOUR", "DECISION TREE CLASSIFIER", "ADABOOST", "BAGGING", "ANN".

APPROACH:

- 1.LOAD THE REQUIRED LIBRARIES SUCH AS PANDAS,MATPLOTLIB,SEABORN ALONG WITH GIVEN DATASET.
- 2.PERFORM EDA ON THE GIVEN DATASET.
- 3.CONVERT ALL THE REQUIRED COLUMNS INTO NUMERIAL COLUMNS USING GET DUMMIES FUNCTION FROM PANDAS LIBRARY.
- 4.CONVERTING ALL REQUIRED FEATURES IN NUMERIAL , CHECK FOR CORRELATION BETWEEN FEATURES AND TARGET AND CONSIDER THE ONLY FEATURES WITH HIGHER CORRELATION.
- 5.IMPORT "LOGISTIC REGRESSION", "SVC", "GAUSSIAN NB", "BERNOULLI NB", "COMPLEMENT NB", "MULTINOMIAL NB", "K NEAREST NEIGHBOUR", "DECISION TREE CLASSIFIER", "ADABOOST", "BAGGING", "ANN", AND SPLIT THE GIVEN DATASET INTO TRAINING AND TESTING DATA USING TRAIN_TEST_SPLIT FUNCTION.THEN CALCULATE ACCURACY SCORE USING SKLEARN LIBRARY BY IMPORTING METRICS.
- 6.ONCE WE GET ACCURACY SCORE OF ALL MODELS FOR BOTH TRAINING AND TESTING DATA, CREATE A DATAFRAME AND LOAD ALL THE ACCURACY OF ALL MODEL.
- 7.VISUALIZATION: ONCE THE DATASET IS CREATED PLOT THE ACCURACIES OF ALL THE MODELS USING BARPLOT USING MATPLOTLIB.

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
pd.set_option("display.Max_columns", 100)
```

C:\Users\Hrishikesh\anaconda3\Lib\site-packages\pandas\core\arrays\masked.py:60: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.5' currently installed).
from pandas.core import (

```
In [2]: data=pd.read_csv(r"C:\Users\Hrishikesh\Desktop\DATA SCIENCE\smart_home_device_usage_data.csv")
data
```

Out[2]:

	UserID	DeviceType	UsageHoursPerDay	EnergyConsumption	UserPreferences	MalfunctionIncidents	DeviceAgeMonths	SmartHomeEfficiency
0	1	Smart Speaker	15.307188	1.961607	1	4	36	1
1	2	Camera	19.973343	8.610689	1	0	29	1
2	3	Security System	18.911535	2.651777	1	0	20	1
3	4	Camera	7.011127	2.341653	0	3	15	0
4	5	Camera	22.610684	4.859069	1	3	36	1
...
5398	5399	Thermostat	4.556314	5.871764	1	0	28	0
5399	5400	Lights	0.561856	1.555992	1	4	24	0
5400	5401	Smart Speaker	11.096236	7.677779	0	0	42	0
5401	5402	Security System	8.782169	7.467929	0	2	28	1
5402	5403	Thermostat	13.540381	9.043076	0	0	30	0

5403 rows × 8 columns



In [3]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5403 entries, 0 to 5402
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   UserID                5403 non-null   int64
1   DeviceType            5403 non-null   object
2   UsageHoursPerDay      5403 non-null   float64
3   EnergyConsumption     5403 non-null   float64
4   UserPreferences       5403 non-null   int64
5   MalfunctionIncidents  5403 non-null   int64
6   DeviceAgeMonths       5403 non-null   int64
7   SmartHomeEfficiency   5403 non-null   int64
dtypes: float64(2), int64(5), object(1)
memory usage: 337.8+ KB
```

In [4]: data.shape

Out[4]: (5403, 8)

In [5]: data.head()

Out[5]:

	UserID	DeviceType	UsageHoursPerDay	EnergyConsumption	UserPreferences	MalfunctionIncidents	DeviceAgeMonths	SmartHomeEfficiency
0	1	Smart Speaker	15.307188	1.961607	1	4	36	1
1	2	Camera	19.973343	8.610689	1	0	29	1
2	3	Security System	18.911535	2.651777	1	0	20	1
3	4	Camera	7.011127	2.341653	0	3	15	0
4	5	Camera	22.610684	4.859069	1	3	36	1

```
In [6]: data.isna().sum()
```

```
Out[6]: UserID                0  
DeviceType                  0  
UsageHoursPerDay            0  
EnergyConsumption           0  
UserPreferences             0  
MalfunctionIncidents        0  
DeviceAgeMonths             0  
SmartHomeEfficiency          0  
dtype: int64
```

```
In [7]: data["DeviceType"].unique()
```

```
Out[7]: array(['Smart Speaker', 'Camera', 'Security System', 'Thermostat',  
              'Lights'], dtype=object)
```

```
In [8]: b=data["DeviceType"].value_counts().reset_index()  
b
```


```
Out[8]:
```

	DeviceType	count
0	Smart Speaker	1108
1	Camera	1101
2	Lights	1087
3	Security System	1068
4	Thermostat	1039

```
In [9]: data.columns
```

```
Out[9]: Index(['UserID', 'DeviceType', 'UsageHoursPerDay', 'EnergyConsumption',  
              'UserPreferences', 'MalfunctionIncidents', 'DeviceAgeMonths',  
              'SmartHomeEfficiency'],  
             dtype='object')
```

```
In [10]: plt.pie(b["count"], labels=b["DeviceType"], autopct="%1.1f%%")  
plt.show()
```



THE ABOVE PIECHART SHOWS THE DISTRIBUTION OF DEVICE TYPES.

```
In [11]: data=data[['DeviceType', 'UsageHoursPerDay', 'EnergyConsumption',  
                  'UserPreferences', 'MalfunctionIncidents', 'DeviceAgeMonths',  
                  'SmartHomeEfficiency']]  
data
```

Out[11]:

	DeviceType	UsageHoursPerDay	EnergyConsumption	UserPreferences	MalfunctionIncidents	DeviceAgeMonths	SmartHomeEfficiency
0	Smart Speaker	15.307188	1.961607	1	4	36	1
1	Camera	19.973343	8.610689	1	0	29	1
2	Security System	18.911535	2.651777	1	0	20	1
3	Camera	7.011127	2.341653	0	3	15	0
4	Camera	22.610684	4.859069	1	3	36	1
...
5398	Thermostat	4.556314	5.871764	1	0	28	0
5399	Lights	0.561856	1.555992	1	4	24	0
5400	Smart Speaker	11.096236	7.677779	0	0	42	0
5401	Security System	8.782169	7.467929	0	2	28	1
5402	Thermostat	13.540381	9.043076	0	0	30	0

5403 rows × 7 columns

```
In [12]: n=pd.get_dummies(data["DeviceType"],drop_first=True).replace({True:1,False:0})
n
```

Out[12]:

	Lights	Security System	Smart Speaker	Thermostat
0	0	0	1	0
1	0	0	0	0
2	0	1	0	0
3	0	0	0	0
4	0	0	0	0
...
5398	0	0	0	1
5399	1	0	0	0
5400	0	0	1	0
5401	0	1	0	0
5402	0	0	0	1

5403 rows × 4 columns

```
In [13]: data=pd.concat([data,n],axis=1).drop(columns="DeviceType")
```

```
In [14]: data=data[['UsageHoursPerDay', 'EnergyConsumption', 'UserPreferences',
                    'MalfunctionIncidents', 'DeviceAgeMonths',
                    'Lights', 'Security System', 'Smart Speaker', 'Thermostat', 'SmartHomeEfficiency']]
```


In [15]: data

Out[15]:

	UsageHoursPerDay	EnergyConsumption	UserPreferences	MalfunctionIncidents	DeviceAgeMonths	Lights	Security System	Smart Speaker	Thermostat	S
0	15.307188	1.961607	1	4	36	0	0	1	0	
1	19.973343	8.610689	1	0	29	0	0	0	0	
2	18.911535	2.651777	1	0	20	0	1	0	0	
3	7.011127	2.341653	0	3	15	0	0	0	0	
4	22.610684	4.859069	1	3	36	0	0	0	0	
...
5398	4.556314	5.871764	1	0	28	0	0	0	1	
5399	0.561856	1.555992	1	4	24	1	0	0	0	
5400	11.096236	7.677779	0	0	42	0	0	1	0	
5401	8.782169	7.467929	0	2	28	0	1	0	0	
5402	13.540381	9.043076	0	0	30	0	0	0	1	

5403 rows × 10 columns




```
In [16]: data.corr()
```

```
Out[16]:
```

	UsageHoursPerDay	EnergyConsumption	UserPreferences	MalfunctionIncidents	DeviceAgeMonths	Lights	Security System	Smart Speaker
UsageHoursPerDay	1.000000	0.026253	0.002161	0.015126	0.025132	-0.000026	-0.002724	-0.005574
EnergyConsumption	0.026253	1.000000	0.016709	-0.000894	-0.016691	-0.001642	0.014465	-0.032159
UserPreferences	0.002161	0.016709	1.000000	0.002154	-0.013781	0.002517	-0.007951	0.007316
MalfunctionIncidents	0.015126	-0.000894	0.002154	1.000000	-0.027023	-0.015969	-0.006193	-0.010185
DeviceAgeMonths	0.025132	-0.016691	-0.013781	-0.027023	1.000000	-0.011805	0.005075	0.008420
Lights	-0.000026	-0.001642	0.002517	-0.015969	-0.011805	1.000000	-0.249095	-0.254896
Security System	-0.002724	0.014465	-0.007951	-0.006193	0.005075	-0.249095	1.000000	-0.252104
Smart Speaker	-0.005574	-0.032159	0.007316	-0.010185	0.008420	-0.254896	-0.252104	1.000000
Thermostat	0.003834	0.015264	-0.009124	0.025397	0.006678	-0.244872	-0.242190	-0.244872
SmartHomeEfficiency	0.183065	-0.178467	0.608713	-0.120836	-0.197077	0.003420	-0.001203	0.003420

```
In [17]: plt.figure(figsize=(18,9))  
sns.heatmap(data.corr(),annot=True)
```

Out[17]: <Axes: >



MACHINE LEARNING:-

```
In [18]: features=data[['UsageHoursPerDay', 'EnergyConsumption', 'UserPreferences',  
                        'MalfunctionIncidents', 'DeviceAgeMonths',  
                        'Lights', 'Security System', 'Smart Speaker', 'Thermostat']]  
target=data['SmartHomeEfficiency']
```

```
In [19]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(features,target)
```

LOGISTIC REGRESSION:-

```
In [20]: from sklearn.linear_model import LogisticRegression  
L=LogisticRegression()  
L.fit(x_train,y_train)
```

```
Out[20]: 

▼ LogisticRegression



LogisticRegression()


```

```
In [21]: L1=L.score(x_train,y_train)*100  
L1
```

```
Out[21]: 87.24086870681145
```

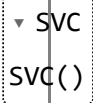
```
In [22]: L2=L.score(x_test,y_test)*100  
L2
```

```
Out[22]: 87.71280532938565
```

```
In [ ]:
```

SUPPORT VECTOR CLASSIFIER:-

```
In [23]: from sklearn.svm import SVC  
S=SVC()  
S.fit(x_train,y_train)
```

```
Out[23]:  SVC()  
SVC()
```

```
In [24]: S1=S.score(x_train,y_train)*100  
S1
```

```
Out[24]: 83.21816386969397
```

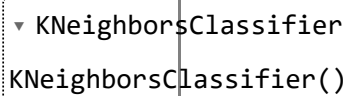
```
In [25]: S2=S.score(x_test,y_test)*100  
S2
```

```
Out[25]: 82.08734270910438
```

```
In [ ]:
```

KNN:-

```
In [26]: from sklearn.neighbors import KNeighborsClassifier  
K=KNeighborsClassifier()  
K.fit(x_train,y_train)
```

```
Out[26]:  KNeighborsClassifier()  
KNeighborsClassifier()
```

```
In [27]: K1=K.score(x_train,y_train)*100  
K1
```

```
Out[27]: 83.24284304047383
```

```
In [28]: K2=K.score(x_test,y_test)*100  
K2
```

```
Out[28]: 71.50259067357513
```

```
In [ ]:
```

NAVIES BAYES:-

```
In [29]: from sklearn.naive_bayes import ComplementNB,BernoulliNB,MultinomialNB,GaussianNB
```

```
In [30]: C=ComplementNB()  
B=BernoulliNB()  
G=GaussianNB()  
M=MultinomialNB()
```

ComplementNB:-

```
In [31]: C.fit(x_train,y_train)
```

```
Out[31]: 

▼ ComplementNB



ComplementNB()


```

```
In [32]: C1=C.score(x_train,y_train)*100  
C1
```

```
Out[32]: 68.06515301085884
```

```
In [33]: C2=C.score(x_test,y_test)*100  
C2
```

```
Out[33]: 67.9496669133975
```

BernoulliNB:-

```
In [34]: B.fit(x_train,y_train)
```

```
Out[34]: ▼ BernoulliNB  
BernoulliNB()
```

```
In [35]: B1=B.score(x_train,y_train)*100  
B1
```

```
Out[35]: 79.14610069101678
```

```
In [36]: B2=B.score(x_test,y_test)*100  
B2
```

```
Out[36]: 79.34863064396744
```

```
In [ ]:
```

GaussianNB:-

```
In [37]: G.fit(x_train,y_train)
```

```
Out[37]: ▼ GaussianNB  
GaussianNB()
```

```
In [38]: G1=G.score(x_train,y_train)*100  
G1
```

```
Out[38]: 83.21816386969397
```

```
In [39]: G2=G.score(x_test,y_test)*100  
G2
```

```
Out[39]: 83.4937083641747
```

```
In [ ]:
```

MultinomialNB:-

```
In [40]: M.fit(x_train,y_train)
```

```
Out[40]: ▼ MultinomialNB  
MultinomialNB()
```

```
In [41]: M1=M.score(x_train,y_train)*100  
M1
```

```
Out[41]: 69.69397828232971
```

```
In [42]: M2=M.score(x_test,y_test)*100  
M2
```

```
Out[42]: 69.94818652849742
```

```
In [ ]:
```


DECISION TREE:-

```
In [43]: from sklearn.tree import DecisionTreeClassifier
D=DecisionTreeClassifier()
D.fit(x_train,y_train)
```

```
Out[43]: ▾ DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [44]: D1=D.score(x_train,y_train)*100
D1
```

```
Out[44]: 100.0
```

```
In [45]: D2=D.score(x_test,y_test)*100
D2
```

```
Out[45]: 91.85788304959289
```

```
In [ ]:
```

ENSEMBLE:-

ADABOAST:-

```
In [46]: from sklearn.ensemble import AdaBoostClassifier, BaggingClassifier, StackingClassifier, RandomForestClassifier
A=AdaBoostClassifier()
A.fit(x_train,y_train)
```

```
Out[46]: ▾ AdaBoostClassifier
AdaBoostClassifier()
```

```
In [47]: A1=A.score(x_train,y_train)*100  
A1
```

```
Out[47]: 94.54590325765054
```

```
In [48]: A2=A.score(x_test,y_test)*100  
A2
```

```
Out[48]: 94.52257586972614
```

BAGGING:-

```
In [49]: b=BaggingClassifier()  
b.fit(x_train,y_train)
```

```
Out[49]: ▼ BaggingClassifier  
BaggingClassifier()
```

```
In [50]: b1=b.score(x_train,y_train)*100  
b1
```

```
Out[50]: 98.81539980256665
```

```
In [51]: b2=b.score(x_test,y_test)*100  
b2
```

```
Out[51]: 94.74463360473723
```


```
In [ ]:
```


DEEP LEARNING:-


```
In [52]: import tensorflow  
from tensorflow import keras
```


```
In [53]: model=keras.Sequential([  
    keras.layers.Dense(8,input_shape=(9,),activation="relu"),  
    keras.layers.Dense(10,activation="relu"),  
    keras.layers.Dense(12,activation="relu"),  
    keras.layers.Dense(1,activation="sigmoid")  
)  
model.compile(optimizer="adam",loss="binary_crossentropy",metrics=["accuracy"])
```


```
In [54]: model.fit(x_train,y_train,epochs=50)
```


Epoch 1/50
127/127  1s 871us/step - accuracy: 0.5907 - loss: 0.7091


Epoch 2/50
127/127  0s 760us/step - accuracy: 0.6562 - loss: 0.6104


Epoch 3/50
127/127  0s 743us/step - accuracy: 0.6625 - loss: 0.6048


Epoch 4/50
127/127  0s 745us/step - accuracy: 0.7099 - loss: 0.5644


Epoch 5/50
127/127  0s 750us/step - accuracy: 0.7318 - loss: 0.5387


Epoch 6/50
127/127  0s 759us/step - accuracy: 0.7974 - loss: 0.4962


Epoch 7/50
127/127  0s 758us/step - accuracy: 0.8324 - loss: 0.4389


Epoch 8/50
127/127  0s 749us/step - accuracy: 0.8582 - loss: 0.3954


Epoch 9/50
127/127  0s 769us/step - accuracy: 0.8689 - loss: 0.3499


Epoch 10/50
127/127  0s 888us/step - accuracy: 0.8731 - loss: 0.3461


Epoch 11/50
127/127  0s 891us/step - accuracy: 0.8647 - loss: 0.3510


Epoch 12/50
127/127  0s 745us/step - accuracy: 0.8647 - loss: 0.3538


Epoch 13/50
127/127  0s 743us/step - accuracy: 0.8551 - loss: 0.3648


Epoch 14/50
127/127  0s 753us/step - accuracy: 0.8806 - loss: 0.3126


Epoch 15/50
127/127  0s 753us/step - accuracy: 0.8732 - loss: 0.3260

Epoch 16/50
127/127  0s 720us/step - accuracy: 0.8648 - loss: 0.3524

Epoch 17/50
127/127  0s 786us/step - accuracy: 0.8717 - loss: 0.3367

Epoch 18/50
127/127  0s 777us/step - accuracy: 0.8655 - loss: 0.3401

Epoch 19/50
127/127  0s 954us/step - accuracy: 0.8754 - loss: 0.3380

Epoch 20/50
127/127  0s 754us/step - accuracy: 0.8695 - loss: 0.3464

Epoch 21/50

127/127  0s 760us/step - accuracy: 0.8706 - loss: 0.3364
Epoch 22/50

127/127  0s 758us/step - accuracy: 0.8793 - loss: 0.3326
Epoch 23/50

127/127  0s 748us/step - accuracy: 0.8640 - loss: 0.3424
Epoch 24/50

127/127  0s 759us/step - accuracy: 0.8824 - loss: 0.3109
Epoch 25/50

127/127  0s 742us/step - accuracy: 0.8697 - loss: 0.3198
Epoch 26/50

127/127  0s 749us/step - accuracy: 0.8812 - loss: 0.3253
Epoch 27/50

127/127  0s 929us/step - accuracy: 0.8675 - loss: 0.3476
Epoch 28/50

127/127  0s 754us/step - accuracy: 0.8745 - loss: 0.3413
Epoch 29/50

127/127  0s 787us/step - accuracy: 0.8768 - loss: 0.3348
Epoch 30/50

127/127  0s 766us/step - accuracy: 0.8814 - loss: 0.3388
Epoch 31/50

127/127  0s 752us/step - accuracy: 0.8779 - loss: 0.3162
Epoch 32/50

127/127  0s 760us/step - accuracy: 0.8776 - loss: 0.3308
Epoch 33/50

127/127  0s 769us/step - accuracy: 0.8851 - loss: 0.3147
Epoch 34/50

127/127  0s 769us/step - accuracy: 0.8802 - loss: 0.3320
Epoch 35/50

127/127  0s 919us/step - accuracy: 0.8735 - loss: 0.3427
Epoch 36/50


127/127  0s 768us/step - accuracy: 0.8866 - loss: 0.3065
Epoch 37/50

127/127  0s 757us/step - accuracy: 0.8808 - loss: 0.3255
Epoch 38/50

127/127  0s 759us/step - accuracy: 0.8834 - loss: 0.3151
Epoch 39/50

127/127  0s 762us/step - accuracy: 0.8747 - loss: 0.3149
Epoch 40/50

127/127  0s 759us/step - accuracy: 0.8792 - loss: 0.3214
Epoch 41/50

127/127  0s 789us/step - accuracy: 0.8769 - loss: 0.3435
Epoch 42/50

```
127/127 ————— 0s 916us/step - accuracy: 0.8839 - loss: 0.3074
Epoch 43/50
127/127 ————— 0s 761us/step - accuracy: 0.8861 - loss: 0.2952
Epoch 44/50
127/127 ————— 0s 766us/step - accuracy: 0.8828 - loss: 0.3259
Epoch 45/50
127/127 ————— 0s 757us/step - accuracy: 0.8773 - loss: 0.3228
Epoch 46/50
127/127 ————— 0s 775us/step - accuracy: 0.8776 - loss: 0.3201
Epoch 47/50
127/127 ————— 0s 773us/step - accuracy: 0.8743 - loss: 0.3317
Epoch 48/50
127/127 ————— 0s 919us/step - accuracy: 0.8881 - loss: 0.3041
Epoch 49/50
127/127 ————— 0s 761us/step - accuracy: 0.8780 - loss: 0.3134
Epoch 50/50
127/127 ————— 0s 729us/step - accuracy: 0.8875 - loss: 0.3107
```

Out[54]: <keras.src.callbacks.history.History at 0x1c028da1f90>

```
In [55]: T=model.predict(x_test)
T
```

```
43/43 ————— 0s 1ms/step
```

Out[55]: array([[0.8194479],
[0.16767004],
[0.50828475],
...,
[0.03387482],
[0.8636771],
[0.51185]], dtype=float32)

```
In [56]: pred=[]
```

```
In [57]: for i in T:
          if i>.5:
              pred.append(1)
          else:
              pred.append(0)
```

```
In [58]: pred
```

```
Out[58]: [1,
           0,
           1,
           1,
           0,
           0,
           0,
           1,
           0,
           1,
           0,
           0,
           0,
           1,
           1,
           0,
           0,
           0,
           0,
           ^
```

```
In [59]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

          accuracy_score(y_test, pred)*100
```

```
Out[59]: 89.04515173945225
```

```
In [60]: confusion_matrix(y_test, pred)
```


```
Out[60]: array([[784,  58],
                [ 90, 419]], dtype=int64)
```



```
In [61]: print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.90	0.93	0.91	842
1	0.88	0.82	0.85	509
accuracy			0.89	1351
macro avg	0.89	0.88	0.88	1351
weighted avg	0.89	0.89	0.89	1351

```
In [62]: train=model.evaluate(x_train,y_train)[1]*100  
train
```

127/127  0s 604us/step - accuracy: 0.8932 - loss: 0.3058

Out[62]: 88.69693875312805

```
In [63]: test=model.evaluate(x_test,y_test)[1]*100  
test
```

43/43  0s 831us/step - accuracy: 0.8906 - loss: 0.3166

Out[63]: 89.04514908790588

```
In [ ]:
```


ACCURACY GRAPH:-

```
In [64]: A={"METHODS":["LOGISTIC REGRESSION","SVC","GAUSSIAN NB","BERNOULLI NB","COMPLEMENT NB","MULTINOMIAL NB","K NEAREST NEI"]
A=pd.DataFrame(A)
A=np.around(A,2)
A
```

Out[64]:

	METHODS	TRAIN ACCURACY	TEST ACCURACY
0	LOGISTIC REGRESSION	87.24	87.71
1	SVC	83.22	82.09
2	GAUSSIAN NB	83.22	83.49
3	BERNOULLI NB	79.15	79.35
4	COMPLEMENT NB	68.07	67.95
5	MULTINOMIAL NB	69.69	69.95
6	K NEAREST NEIGHBOUR	83.24	71.50
7	Decision Tree Classifier	100.00	91.86
8	ADABOOST	94.55	94.52
9	BAGGING	98.82	94.74
10	ANN	88.70	89.05

```
In [65]: plt.bar(A["METHODS"],A["TRAIN ACCURACY"],width=0.3,label="TRAINING ACCURACY")
plt.bar(A["METHODS"],A["TEST ACCURACY"],align="edge",width=0.3,label="TESTING ACCURACY")
plt.legend(bbox_to_anchor=[1,0,0,1])
plt.xlabel("METHODS----->")
plt.ylabel("ACCURACY----->")
plt.xticks(rotation=90)
plt.show()
```



CONCLUSION:

FROM THE ABOVE BAR CHART IT IS CLEAR THAT LOGISTIC REGRESSION AND ADABOAST ARE BEST

In []: